
Japanese Digit Classification Using Convolutional Neural Networks

Max Ardito

maxwell.ardito@mail.mcgill.ca

Edwin Meriaux

edwin.meriaux@mail.mcgill.ca

Ohood Sabr

ohood.sabr@mail.mcgill.ca

Abstract

In this paper, the performance of several Convolutional Neural Network (CNN) designs is evaluated using a variation on the MNIST dataset that includes both western Arabic numerals and Japanese numerals. Each image contains a single Japanese numeral and may also potentially have up to three western numerals. The goal of the classifier is to successfully classify the Japanese digit. This study examined the effectiveness and precision of both a base CNN designs and a CNN design that includes residual layers, as well as squeeze-and-excitation layers. The model that scored the highest accuracy is the later model, a variation on the famous ResNet-18 classifier. Preprocessing methods including digit rotation are also discussed.

Keywords: Convolutional neural network, CNN, Image classification, MNIST dataset, Residual layers.

1 Introduction

Computer vision (CV) deals with the issues and strategies of simulating human vision using computers. Its primary function is to analyze and comprehend images and videos in order to make conclusions about their content. Image recognition is a type of technology that employs machine learning models to analyze, evaluate, and classify various entities within images[1].

The CV field has greatly evolved during the last few decades. Convolution Neural Networks (CNNs) are a well-known CV model[2]. The standard CNN is a deep neural network made up of various layers of neurons that perform tasks such as filtering, pooling, and activation. Filtering and pooling layers in carry out feature extraction, while a fully connected layer perform feature mapping to produce the final output after being passed through activation functions[3]. What distinguishes a CNN from a "simple" multilayer perceptron (MLP) network is therefore this combined use of convolutional filtering layers, pooling, and non-linearities like tanh, sigmoid, and ReLU functions [4].

The focus of this work is to examine the classification performance of the CNN model in given a variation on the notable MNIST dataset that contains a Japanese numeral along with up to three western Arabic numerals in a single image. The class labels of these images refer to the digit value of the Japanese numeral.

This report is organized as follows: Section 2 discusses the data set and talks about the preprocessing steps that were taken. Section 3 presents the various models experimented with during the training phase. Section 4 presents the findings, and Section 5 concludes with some possible further research.

2 Datasets

2.1 Overview

The given dataset is a variation on the the MNIST data set which is a standard dataset given by the Modified National Institute of Standards and Technology [5]. The traditional dataset contains images, each of a handwritten digit (0-9) as well as their class label values. The variation of MNIST given for this project is a set of images, each of which contain a single handwritten Japanese digits and up to three western digits. The overall dataset consists of 60,000 training samples (images). The objective is to identify the Japanese digit present in the input image. Using CNNs, the multi-class image classification issue is resolved. After extracting images from the dataset, we have seen that each image has 28-pixel width and a 28-pixel height (28 X 28) and a grayscale (single channel) image. The intensity values for each pixel varies from 0 (black) to 255 (white). Figure 1 displays an example image from the dataset in question.




Figure 1: Traing image sample.

2.2 Preprocessing

Before the training data is fed into the different CNN designs, it goes through a number of preprocessing stages. Initially, the images are transformed from Numpy arrays into Pytorch tensors to enable high dimensional arithmetic operations on the GPU. In the next stage, the images are normalized so that their pixel values fall between -1 and 1. The advantage of normalizing the input data is that it prevents the training process from being difficult due to huge gradient values[5]. One other method that was tried was Sobel filtering to process images, however, this did not work out well mainly due to the fact that the images are only 28x28 pixels. The Sobel filter just created an output image that was effectively the same image due to the MNIST's small image size. The CNN models can also be fed directly with raw values using the greyscale pixels, however, doing so may result in longer training times and lower accuracy.

Aside from image quality improvement, augmentation was also used. This is used in order to expand the number of training instances without using new images, augmenting the original images and using both the original and augmented images for training. Another method is image rotation. As the name implies it means the images used in training set are rotated. This is done to augment the training dataset. This increases variation from the images in the training phase. Additionally, the handwritten digits in the MNIST dataset are not aligned at the same angle. This means that training the algorithm on images that are rotated should, in theory, also increase the accuracy. However, image rotation benefits were not seen in the output, as the accuracy in training and validation actually decreased. A possible way to improve this in the future would be to only rotate individual digits instead of the entire image. This can be done by separating the characters by recognizing the outline of each individual number before rotating and then merging the values into one image. A final preprocessing method used was image color augmentation. This works on modifying the pixel colors and intensity. This also did not improve accuracy, in part because the images are not in color.




Figure 2: NN (Left) and CNN (Right) comparison.

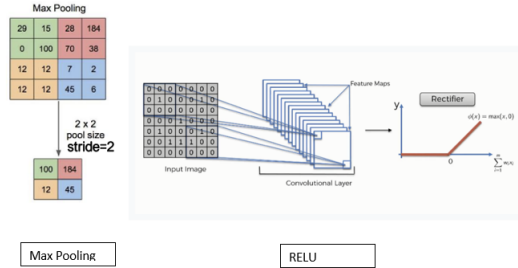


Figure 3: Visualizations of Max Pooling (Left) and ReLU activation (Right).

3 Proposed Approach

3.1 Base CNN Model

A Convolutional Neural Network (CNN) is a form of Neural Networks (NN) (Figure: 2) that works to try and reduce the computational complexity of a NN by attaching neurons on one layer only to a select number of neurons from the previous layer. This reduces the computation by reducing the number of features that each neurons has to take in. This makes the computation on each node lower and the back propagation faster as well. As the layers move from the broadest to the end the number of neurons generally decrease to an eventual single end neurons that takes input from all the previous neurons.

CNNs also contain additional layers for averaging and activation. For instance, Rectified Linear Activation Unit (RELU) layers are used to further help reduce computational complexity (Figure: 3). Effectively this is just a threshold value. In a one dimensional input case if the value is above a certain value then the output is a linear value greater than 0 but if it bellows a certain value then the output is 0. This threshold value is calculated as part of the training of the system to determine if the value is useful in the following layers. Turning the value to 0 makes it so that the next layer does not care for the output of this neuron.

Another layer used in CNNs to reduce the complexity is the pooling layers (Figure: 3). This filter reduces an image to a fraction of its current size by downsampling. It breaks up an image into specific rectangles and the output of the filter has each of those rectangles as just one pixel in the new image. This is done by averaging all the values the rectangle contained.

Various other reduction layers can be implemented in the base CNN model as well. The flatten function takes a PyTorch tensor and reshapes it into a one-dimensional tensor. This is similar to the squeeze function in NumPy which takes a NumPy matrix and turns it into a one dimension array. Adaptive Average Pool layer (Adaptiveavgpool) is a special pooling function that does not takes data from just one plane but multiple. The planes can be different inputs into the neutron in question. The linear activation layer is a different activation function. It is based off of linear regression and trained on the input values from the previous layer's neutrons.

3.2 Residual Layers

The aforementioned CNN layers are standard to all CNNs. These base layers can be configured in many different ways and their order and parameters resemble hyperparameters in the training process. Based on analyzing literature pertaining to models with high accuracy metrics on the original MNIST dataset, two additional layers were used in training a variation on the base CNN model.

The first type of additional CNN layer used is the residual layer, which allows for a sort of regularization in the network that solves the issue of vanishing or exploding gradient. Layers can be skipped and thus connected to the output of previous layers in the network in a non-sequential manner. The effect that this has on the network is not unlike that of LSTM blocks on recurrent neural networks (RNN). A semblance of memory is attained in the network and this regularization has been proven to increase accuracy on image recognition problems [6].

3.3 Squeeze and Excitation Layers

The concept of retaining both a balance of high and low-level features is perhaps the central conceit of CNN architectures. Learning adequate weights for high-level neurons that classify large objects and shapes is just as important as learning the weights for low-level neurons that distinguish edges and contours [7]. While residual layers attempt to establish this balance, squeeze and excitation (SE) layers take this one step further.

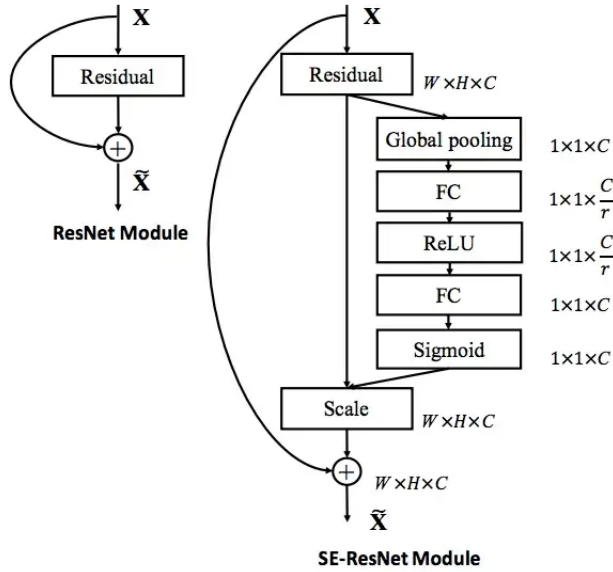


Figure 4: Squeeze-and-Excitation Layer (Right) compared to Residual Layer (Left) [7]

An SE layer is very similar to a residual layer, except in this case each channel is pooled or "squeezed" into a single value, passed to a ReLU activation function, reduced by a certain ratio r , and activated once more by the sigmoid function. Figure 4 shows a full diagram of the SE architecture compared to a residual layer.

4 Results

Based on techniques discussed in the previous sections, we constructed three different CNN architectures. The first architecture was a simple base CNN with no preprocessing and no residual layers. The second CNN was a model inspired roughly by the ResNet-18 model, with various changes including the addition of residual layers and SE layers. The third CNN used was simply the same as the second CNN model but with additional preprocessing steps. Each model was tested on a

validation set consisting of 0.1% of the original training dataset (6,000 images). A batch size of 128 images was used on all models, and the number of epochs used to train each model varied between 2-15 depending on the accuracy metric of the first epoch.

The base model stood as a sort of benchmark attempt at designing a classifier that functioned properly. The CNN architecture used here consisted of two convolutional blocks, each of which consisted of a fully connected layer, a normalization layer, a ReLU layer, and a max pooling layer. Blocks were output to a linear layer which mapped the final number of output channels to 10 neurons, the maximum of which being the classified digit. Kernel sizes of 3 pixels were used on the image filters along with strides and paddings of 1 pixel. The base model scored 91% accuracy on the validation set but only 85% on the testing set.

Once this benchmark was established, a blueprint of ResNet-18 was used as a starting point to construct a network that yielded the highest accuracy. For this model, four convolutional blocks were constructed that consisted of a fully connected convolutional layer, a residual layer whose channels are sometimes downsampled, an SE layer with an r value that peaks in the middle of the network, ReLU activation, and average pooling. The final block contains no convolutional layers, but instead two residual layers that channels, average, flatten, and output to the 10 classifying neurons. This model scored 99% on validation data and 92% on the testing data.

Finally, rotational preprocessing was performed on the dataset, straightening out angles in the digits. While this improved accuracy on some training phases, the difference deemed the preprocessing step unnecessary.

5 Future Work (K-Means) and Conclusion

While the ResNet variation model performed quite well on the dataset, more refined preprocessing steps could indeed be taken given more time. For instance, each image could potentially be sliced up into four quadrants, and given the prior knowledge that each image contains up to four digits, one Japanese digit, and up to three western digits, a pre-classifier could be constructed using Bayesian probability and K-means clustering to classify which digit is the Japanese digit. It is estimated that this would greatly increase training, since passing a single Japanese digit into a CNN and estimating its value is a much simpler task than passing four images.

In order to do this, each image would be sliced up and a K-means clustering algorithm would be used to learn the differences between the images. This learned difference would distinguish blank spaces, western Digits, and Japanese digits. Then, using this data along with each type of digit's probability of occurring in a given quadrant, the model would identify the Japanese digit in each image and then pass it through the digit classifying CNN.

Additionally, future work could also include principal component analysis (PCA) as a pre-processing step to compress images through the use of an orthogonal transformation. PCA could help narrow down a complex data set's dimensions, and it is clearly stated in various pieces of literature that PCA improves the accuracy of CNNs [8]. Another possible preprocessing method that could also be done is cropping to reduce the amount of blank space in the images. This would reduce training on the blank spaces and only use the minimum amount of pixels of interest. Each image would be broken down into rectangles that act as bounding boxes for the characters, processing only the pixels in the boxes. A final possibility could be performing image segmentation. This breaks images down into multiple sections, effectively partitioning image content down to their boundaries. If this is done then by knowing the colors of each section a threshold can be used to know which partitions are characters. This can be used in the K-Means method or any method that requires determining the specific number and no extra pixels. This is also would by itself change the colors in the space by pushing the segments of interest to pure white colors and the rest to pure black. This would remove noise in the data associated with variations in the greyscale.

6 Statement of Contribution

1. Max Ardito: Report editing, Residual Layers, SE Layers, Training Architecture
2. Ohood Sabr: Report writing, Preprocessing (rotational method)
3. Edwin Meriaux: Report writing, Preprocessing (rotation and image augmentation), Basic RELU CNN algorithm, training and testing

References

- [1] A. A. M. Al-Saffar, H. Tao, and M. A. Talab, "Review of deep convolution neural network in image classification," in *2017 International Conference on Radar, Antenna, Microwave, Electronics, and Telecommunications (ICRAMET)*, pp. 26–31, 2017.
- [2] W. Rawat and Z. Wang, "Deep convolutional neural networks for image classification: A comprehensive review," *Neural computation*, vol. 29, no. 9, pp. 2352–2449, 2017.
- [3] R. Yamashita, M. Nishio, R. K. G. Do, and K. Togashi, "Convolutional neural networks: an overview and application in radiology," *Insights into imaging*, vol. 9, no. 4, pp. 611–629, 2018.
- [4] A. F. Agarap, "An architecture combining convolutional neural network (cnn) and support vector machine (svm) for image classification," *arXiv preprint arXiv:1712.03541*, 2017.
- [5] R. Pramoditha, "Acquire, understand and prepare the mnist dataset," 2022–06–04.
- [6] N. Armanfard, "Ecse 551 - machine learning for engineers lecture 21 — recurrent neural networks (rnns)," December 2022.
- [7] P.-L. Pröve, "Squeeze-and-excitation networks," 2017-11-07.
- [8] Y. K. Taehong Kwak, Ahram Song, "The effect of the principal component analysis in convolutional neural network for hyperspectral image classification," *The 40th Asian Conference on Remote Sensing (ACRS 2019)*, 2019.

7 Appendix A: Code

```
1  # -*- coding: utf-8 -*-
2  """Squeeze_And_Excitation.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1pe0HRavaQ9INTuXjxKeM4kah2NzLRute
8  """
9
10 # Code Citations:
11 # https://medium.com/@krishna.ramesh.tx/training-a-cnn-to-distinguish-between-
12 # mnist-digits-using-pytorch-620f06aa9ffa
13 # https://www.analyticsvidhya.com/blog/2019/10/building-image
14 # -classification-models-cnn-pytorch/
15
16 """## Import Packages"""
17
18 # Commented out IPython magic to ensure Python compatibility.
19 import pickle
20 # importing the libraries
21 import pandas as pd
22 import numpy as np
23 import cv2
24 from google.colab.patches import cv2_imshow
25 # for reading and displaying images
26 from skimage.io import imread
```

```

27 import matplotlib.pyplot as plt
28 # %matplotlib inline
29
30 # for creating validation set
31 from sklearn.model_selection import train_test_split
32
33 # for evaluating the model
34 from sklearn.metrics import accuracy_score
35 from tqdm import tqdm
36
37 # PyTorch libraries and modules
38 import torch
39 from torch.autograd import Variable
40 from torch.nn import Linear, ReLU, CrossEntropyLoss, Sequential, Conv2d,
    ↳ MaxPool2d, Module, Softmax, BatchNorm2d, Dropout
41 from torch.optim import Adam, SGD
42
43 # Commented out IPython magic to ensure Python compatibility.
44 import pickle
45 # importing the libraries
46 import pandas as pd
47 import numpy as np
48 import cv2
49 from google.colab.patches import cv2_imshow
50 # for reading and displaying images
51 from skimage.io import imread
52 import matplotlib.pyplot as plt
53 # %matplotlib inline
54
55 # for creating validation set
56 from sklearn.model_selection import train_test_split
57
58 # for evaluating the model
59 from sklearn.metrics import accuracy_score
60 from tqdm import tqdm
61
62 # PyTorch libraries and modules
63 import torch
64 from torch.autograd import Variable
65 from torch.nn import Linear, ReLU, CrossEntropyLoss, Sequential, Conv2d,
    ↳ MaxPool2d, Module, Softmax, BatchNorm2d, Dropout
66 from torch.optim import Adam, SGD
67 #pytorch utility imports
68 import torch
69 import torchvision
70 import torchvision.transforms as transforms
71 from torch.utils.data import DataLoader, TensorDataset
72 from torchvision.utils import make_grid
73 import pickle
74 #neural net imports
75 import torch.nn as nn
76 import torch.nn.functional as F
77 import torch.optim as optim
78 from torch.autograd import Variable
79 #import external libraries
80 import pandas as pd
81 import numpy as np
82 import matplotlib.pyplot as plt
83 from sklearn.model_selection import train_test_split
84 import os
85 import math
86 # %matplotlib inline
87
88 """## Mount Google Drive

```

```

89
90 Make sure you have all of the datasets in your google drive.
91 """
92
93 from google.colab import drive
94 drive.mount('/content/drive')
95
96 """_NOTE: In the code block below, make sure you change the paths accordingly
97 ↪ to your drive :_)_"""
98
99 # Training data and labels
100 data_x = pickle.load( open('drive/MyDrive/Train.pkl', 'rb'))
101 data_y = np.genfromtxt('drive/MyDrive/Train_labels.csv', delimiter=',')
102
103 # Test data for Kaggle submission
104 test_x = pickle.load( open( 'drive/MyDrive/Test.pkl', 'rb' ))
105 test_y = np.genfromtxt('drive/MyDrive/ExampleSubmissionRandom.csv',
106 ↪ delimiter=',')
107
108 # Sample a random image from the dataset and print its class label
109 test_sample = 155
110 plt.imshow(data_x[test_sample][0], cmap='gray', interpolation="bicubic")
111 plt.show()
112 print("Label: ", data_y[test_sample][1])
113 print(data_x.shape[0])
114
115 # Image Rotation - NOT WORKING
116 # from keras_preprocessing.image import ImageDataGenerator, array_to_img,
117 ↪ img_to_array, load_img
118 # #img = load_img('Screenshot_47.png')
119
120 # from numpy import expand_dims
121 # from keras_preprocessing.image import load_img
122 # from keras_preprocessing.image import img_to_array
123 # from keras_preprocessing.image import ImageDataGenerator
124 # import matplotlib.pyplot as plt
125
126 # rotated_train_images = []
127 # rotated_train_labels = []
128
129 # for j in range(train_images.shape[0]):
130 # #for j in range(10):
131 # # load the image
132 # plt.figure(figsize=(45,30))
133
134 # # convert to numpy array
135 # data = img_to_array(train_images[j])
136 # rotated_train_images.append(train_images[j])
137 # rotated_train_labels.append(train_labels[j])
138 # #print(data.shape)
139 # # expand dimension to one sample
140 # samples = expand_dims(data, 0)
141
142 # datagen = ImageDataGenerator(featurewise_center=True,
143 # rotation_range=(0-10),brightness_range=[0,1])
144 # # prepare iterator
145 # it = datagen.flow(samples, batch_size=1)
146
147 # # generate samples and plot
148 # for i in range(6):
149 # # define subplot
150 # #plt.subplot(330 + 1 + i)
151 # # generate batch of images
152 # batch = it.next()

```



```

150 #         # convert to unsigned integers for viewing
151 #         image = batch[0].astype('uint8')
152 #         rotated_train_images.append(image)
153 #         # plot raw pixel data
154 #         rotated_train_labels.append(train_labels[j])
155 #     print(len(rotated_train_images))
156 #     print(len(rotated_train_labels))
157
158 """## Split and Convert Data To Pytorch Tensors"""
159
160 data_x_tmp = []
161 test_x_tmp = []
162 for i in data_x:
163     data_x_tmp.append(i[0])
164
165 for i in test_x:
166     test_x_tmp.append(i[0])
167
168 data_x = np.asarray(data_x_tmp)
169 test_x = np.asarray(test_x_tmp)
170 print("Training Data Shape: ", data_x.shape)
171 print("Test Data Shape: ", test_x.shape)
172
173 # Remove invalid first element and index elements from class labels
174 data_y = data_y[1:]
175 data_y = data_y[:,1]
176 data_y = np.int_(data_y)
177
178 test_y = test_y[1:]
179 test_y = test_y[:,1]
180 test_y = np.int_(test_y)
181
182 train_x, val_x, train_y, val_y = train_test_split(data_x, data_y, test_size =
↪ 0.1)
183 (train_x.shape, train_y.shape), (val_x.shape, val_y.shape)
184
185 # Convert training images into torch format
186 train_x = train_x.reshape(54000, 1, 28, 28)
187 train_x = torch.from_numpy(train_x)
188
189 test_x = test_x.reshape(10000, 1, 28, 28)
190 test_x = torch.from_numpy(test_x)
191
192 # converting the target into torch format
193 train_y = train_y.astype(int);
194 train_y = torch.from_numpy(train_y)
195
196 test_y = test_y.astype(int);
197 test_y = torch.from_numpy(test_y)
198
199 # shape of training data
200 print("Shape of training data and labels")
201 train_x.shape, train_y.shape
202
203 # shape of training data
204 print("Shape of test data")
205 test_x.shape, test_y.shape
206
207 # converting validation images into torch format
208 val_x = val_x.reshape(6000, 1, 28, 28)
209 val_x = torch.from_numpy(val_x)
210
211 # converting the target into torch format
212 val_y = val_y.astype(int);

```

```

213 val_y = torch.from_numpy(val_y)
214
215 # shape of validation data
216 val_x.shape, val_y.shape
217
218 # converting original dataset to final training set
219 final_x = data_x.reshape(60000, 1, 28, 28)
220 final_x = torch.from_numpy(final_x)
221
222 final_y = data_y.astype(int);
223 final_y = torch.from_numpy(final_y)
224
225 batch_size = 128
226
227 # converting training images into torch format
228 train_tensor_x = torch.Tensor(train_x)
229 val_tensor_x = torch.Tensor(val_x)
230 test_tensor_x = torch.Tensor(test_x)
231 final_tensor_x = torch.Tensor(final_x)
232
233 # converting the target into torch format (adding indices)
234 train_tensor_y = torch.Tensor(train_y)
235 val_tensor_y = torch.Tensor(val_y)
236 test_tensor_y = torch.Tensor(test_y)
237 final_tensor_y = torch.Tensor(final_y)
238
239 train_dataset = TensorDataset(train_tensor_x, train_tensor_y) # create your
↳ dataset
240 val_dataset = TensorDataset(val_tensor_x, val_tensor_y) # create your dataset
241 test_dataset = TensorDataset(test_tensor_x, test_tensor_y) # create your dataset
242 final_dataset = TensorDataset(final_tensor_x, final_tensor_y) # create your
↳ dataset
243
244 # Final datasets
245 train_dl = DataLoader(train_dataset, batch_size=batch_size)
246 val_dl = DataLoader(val_dataset, batch_size=batch_size)
247 test_dl = DataLoader(test_dataset, batch_size=batch_size)
248 final_dl = DataLoader(final_dataset, batch_size=batch_size)
249
250 # Display the first N batches of data
251 N = 4
252 for i, (images, labels) in enumerate(train_dl):
253     print('Batch index: ', i)
254     print('Batch size: ', images.size())
255     print('Batch label: ', labels)
256     if(i > N):
257         break
258
259 """## Define The Model"""
260
261 class ResBlock(nn.Module):
262     def __init__(self, in_channels, out_channels, downsample):
263         super().__init__()
264         if downsample:
265             self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
↳ stride=2, padding=1)
266             self.shortcut = nn.Sequential(
267                 nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=2),
268                 nn.BatchNorm2d(out_channels)
269             )
270         else:
271             self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
↳ stride=1, padding=1)
272             self.shortcut = nn.Sequential()

```

```

273
274     self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,
    ↪ stride=1, padding=1)
275     self.bn1 = nn.BatchNorm2d(out_channels)
276     self.bn2 = nn.BatchNorm2d(out_channels)
277
278     def forward(self, input):
279         shortcut = self.shortcut(input)
280         input = nn.ReLU()(self.bn1(self.conv1(input)))
281         input = nn.ReLU()(self.bn2(self.conv2(input)))
282         input = input + shortcut
283         return nn.ReLU()(input)
284
285     # Squeeze and Excite Layer
286     # Code citation:
    ↪ https://towardsdatascience.com/introduction-to-squeeze-excitation-
    ↪ networks-f22ce3a43348
287
288     class SELayer(nn.Module):
289         def __init__(self, channel, reduction=16):
290             super(SELayer, self).__init__()
291             self.avg_pool = nn.AdaptiveAvgPool2d(1)
292             self.fc = nn.Sequential(
293                 nn.Linear(channel, channel // reduction, bias=False),
294                 nn.ReLU(inplace=True),
295                 nn.Linear(channel // reduction, channel, bias=False),
296                 nn.Sigmoid()
297             )
298         def forward(self, x):
299             b, c, _, _ = x.size()
300             y = self.avg_pool(x).view(b, c)
301             y = self.fc(y).view(b, c, 1, 1)
302             return x * y.expand_as(x)
303
304     """
305     CURRENT BEST MODEL
306     """
307     class Net(nn.Module):
308         def __init__(self, resblock, outputs=10):
309             super().__init__()
310             self.network = nn.Sequential(
311                 nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=2),
312                 resblock(32, 32, downsample=False),
313                 SELayer(32, 2),
314                 nn.ReLU(),
315                 nn.AvgPool2d(1, stride=1),
316
317                 nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=2),
318                 resblock(64, 64, downsample=False),
319                 SELayer(64, 16),
320                 nn.ReLU(),
321                 nn.AvgPool2d(1, stride=1),
322
323                 nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=0),
324                 resblock(128, 128, downsample=False),
325                 SELayer(128, 64),
326                 nn.ReLU(),
327                 nn.AvgPool2d(2, stride=2),
328
329                 nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=0),
330                 resblock(256, 256, downsample=False),
331                 resblock(256, 512, downsample=True),
332                 SELayer(512, 2),
333                 nn.ReLU(),
334                 nn.AvgPool2d(2, stride=2),

```

```

335         resblock(512, 512, downsample=False),
336         resblock(512, 1024, downsample=True),
337         nn.AdaptiveAvgPool2d(1),
338         nn.Flatten(),
339         nn.Linear(1024, outputs)
340     )
341
342
343     def forward(self, input):
344         y = self.network(input)
345         return y
346
347     """# Define Training Functions"""
348
349     # Validation Function
350     # Code Citation:
351     ↪ https://medium.com/@krishna.ramesh.tx/training-a-cnn-to-distinguish-between-
352     # -mnist-digits-using-pytorch-620f06aa9ffa
353     def validate(model, data):
354         total = 0
355         correct = 0
356         for i, (images, labels) in enumerate(data):
357             images = images.cuda()
358             x = model(images)
359             value, pred = torch.max(x, 1)
360             pred = pred.data.cpu()
361             total += x.size(0)
362             correct += torch.sum(pred == labels)
363         return correct*100./total
364
365     # Training Function
366     # Code Citation:
367     ↪ https://medium.com/@krishna.ramesh.tx/training-a-cnn-to-distinguish-between-
368     # mnist-digits-using-pytorch-620f06aa9ffa
369     import copy
370     from torchsummary import summary
371
372     def train(num_epochs=3, dataset=train_dl, lr=1e-3, device="cpu"):
373         resnet18 = Net(ResBlock, outputs=10)
374         cnn = resnet18.to(torch.device("cuda:0" if torch.cuda.is_available() else
375         ↪ "cpu"))
376         accuracies = []
377         # cnn = create_model().to(device)
378         cec = nn.CrossEntropyLoss()
379         optimizer = optim.Adam(cnn.parameters(), lr=lr)
380         max_accuracy = 0
381         for epoch in range(num_epochs):
382             for i, (images, labels) in enumerate(dataset):
383                 images = images.to(device)
384                 # labels = labels.to(device)
385                 labels = labels.type(torch.cuda.LongTensor)
386                 optimizer.zero_grad()
387                 pred = cnn(images)
388                 loss = cec(pred, labels)
389                 loss.backward()
390                 optimizer.step()
391             accuracy = float(validate(cnn, val_dl))
392             accuracies.append(accuracy)
393             if accuracy > max_accuracy:
394                 best_model = copy.deepcopy(cnn)
395                 max_accuracy = accuracy
396                 print("Saving Best Model with Accuracy: ", accuracy)
397                 print('Epoch:', epoch+1, "Accuracy :", accuracy, '%')
398         plt.plot(accuracies)

```

```

396     return best_model
397
398     """## Train on Training Data"""
399
400     if torch.cuda.is_available():
401         device = torch.device("cuda:0")
402     else:
403         device = torch.device("cpu")
404     print("No Cuda Available")
405
406     device
407
408     cnn = train(5, dataset=train_dl, device=device)
409
410     """## Train on Entire Dataset"""
411
412     # Train the final model
413     cnn = train(2, dataset=final_dl, device=device)
414
415     """## Predict Testing Set"""
416
417     def predict_final(model, data):
418         y_pred = []
419         for i, (images, labels) in enumerate(data):
420             images = images.cuda()
421             x = model(images)
422             value, pred = torch.max(x, 1)
423             pred = pred.data.cpu()
424             y_pred.extend(list(pred.numpy()))
425         return np.array(y_pred)
426
427     y_pred = predict_final(cnn, test_dl)
428
429     """## Save Testing Set Predictions (CSV)"""
430
431     y_final = np.zeros((len(y_pred), 2))
432     for i in range(len(y_pred)):
433         y_final[i, :] = np.array([i, int(y_pred[i])])
434
435     # save array into csv file
436     np.savetxt("final.csv",
437               y_final,
438               '%i,%i',
439               delimiter = ",")
440

```